

# **System and Device Programming - Lab 5**

**Professor P. Laface**

Tamer Saadeh (222201) <tamer@tamersaadeh.com>

Politecnico di Torino - 28<sup>th</sup> April, 2015

# **Part I.**

## **pack.sh**

The following contains a commented version of the pack.sh script, which is a script that helps with compiling and running a custom UNIX-like OS.

```
1  #!/bin/bash
2
3  # check if parameters provided
4  if [ $# -ne 1 ]
5  then
6  echo "Usage: $0  number of k-bytes"
7  exit 1
8  fi
9
10 kbytes=$1
11 # clean log files
12 rm pack1.txt 1>/dev/null 2>/dev/null
13 rm pack2.txt 1>/dev/null 2>/dev/null
14
15 # save these to chmod files back
16 var_user=laface
17 var_group=laface
18
19 # Creates an empty file of size $kbytes*1024*512 with filled with zeros
20 # of: output file / if: input file / bs: block size / count: how many
21 # blocks / seek: skips forward into file
22 # The file created will have the following structure:
23 # [000.....000][000...000]
24 # ~~~~~
25 # $kbytes*1024*512      512      bytes
26 dd if=/dev/zero of=hd.img bs=512 count=1 seek=$((kbytes*1024))
27
28 # mounts the image as a loop device in /dev/loop1
29 losetup /dev/loop1 hd.img
30
31 # The following will generate a file fdisk.input so that we can script
32 # the formatting of disk.
33 # The meaning of the content of the file:
34 # x: use "expert" mode
35 # h: set the number of heads to 16
36 # s: set the number of sectors/tracks in the disk to 63
37 # c: set the number of cylinders to $kbytes
38 # r: return to main menu, ie exit "expert" mode
39 # n: creates a new physical partition numbered 1
40 # The next two lines use the default values for the new partion for the
41 # first cylinder location (ie 1) and the last cylinder (ie 256)
42 # a: set boot flag for the first partition
43 # w: write modifications to disk file (ie /dev/loop1 or hd.img)
44
45 cat <<EOF > fdisk.input
46 x
47 h
48 16
49 s
50 63
```

```

51 c
52 EOF
53
54 echo $kbytes >> fdisk.input
55
56 cat <<EOF >> fdisk.input
57 r
58 n
59 p
60 1
61
62
63 a
64 1
65 w
66 EOF
67
68 # actually run fdisk with the previous commands
69 fdisk hd.img < fdisk.input 1>>pack1.txt 2>>pack2.txt
70
71 # list partition tables in the disk we just created along with the size
72 # in sectors rather than cylinders
73 fdisk -l -u /dev/loop1 1>>pack1.txt 2>>pack2.txt
74
75 # extract the number of the blocks from the fdisk command
76 # tail -1 : keeps the last line
77 # tr -s " " : squeezes the spaces
78 # cut -d " " -f 5 : splits by space and gets the 5th field, ie the blocks
79 # number with the "+" sign at the end
80 # cut -d "+" -f 1 : get rid of the "+" sign so that we have a pure number
81 blocks=$(fdisk -u -l /dev/loop1 | tail -1 | tr -s " " | \
82 cut -d " " -f 5 | cut -d "+" -f 1)
83
84 # detach the loop-ed disk image
85 losetup -d /dev/loop1
86
87 # remount the disk image with an offset as 63*512 so to skip the MBR
88 # data since these data are stored in the first track and limiting the
89 # size of the disk to the actual size (ie $blocks*1024). This is needed
90 # so that we know where to stop when creating the filesystem.
91 losetup -o $((63*512)) --sizelimit $(($blocks*1024)) /dev/loop0 hd.img \
92 1>>pack1.txt 2>>pack2.txt
93
94 # creates an ext2 filesystem on the disk
95 mkfs.ext2 /dev/loop0
96
97 # create the grub boot menu item
98 # we set the title of the grub menu item to "minimal linux like kernel"
99 # we tell grub where to find the kernel we want to start and on which disk.
100 # Mount it as read-only, be quiet and display a splash screen for the kernel
101 # (if available). Finally, be quiet and don't print anything to the screen
102 # while loading the kernel.
103 cat <<EOF > menu.lst

```

```

104 title minimal linux like kernel
105 kernel /boot/minimal_linux_like_kernel root=/dev/hda ro quiet splash
106 quiet
107
108 EOF
109
110 # mount /dev/loop0 to mount_point
111 mkdir mount_point 1>>pack1.txt 2>>pack2.txt
112 mount /dev/loop0 mount_point 1>>pack1.txt 2>>pack2.txt
113
114 # copy grub files from local hard drive
115 mkdir -p mount_point/boot/grub
116 cp ../grub/stage1 ../grub/stage2 menu.lst mount_point/boot/grub \
117     1>>pack1.txt 2>>pack2.txt
118
119 # copy kernel
120 cp -v src/kernel mount_point/boot/minimal_linux_like_kernel \
121     1>>pack1.txt 2>>pack2.txt
122
123 # umount mount_point
124 umount mount_point 1>>pack1.txt 2>>pack2.txt
125 rm -r mount_point 1>>pack1.txt 2>>pack2.txt
126
127 # detach /dev/loop0
128 losetup -d /dev/loop0 1>>pack1.txt 2>>pack2.txt
129
130 # creates the grub.input file for grub
131 # The following runs in GRUB shell, which means:
132 # device: specifies the device and the image to install grub on
133 # geometry: prints information about the device we specified with
134 #           the given cylinders, head and sector sizes
135 # root: sets the first partition as the root partition for the 2nd
136 #       stage of booting
137 # setup: installs grub to the MBR of the device
138 cat <<EOF > grub.input
139 device (hd0,0) hd.img
140 EOF
141 echo geometry \ (hd0\ ) $kbytes 16 63 >> grub.input
142 cat <<EOF >> grub.input
143 root (hd0,0)
144 setup (hd0)
145 quit
146 EOF
147
148 ../grub/grub.bin --device-map=/dev/null < grub.input \
149     1>>pack1.txt 2>>pack2.txt
150
151 # Set permissions to our user to the logs and image
152 chown ${var_user}:${var_group} pack1.txt
153 chown ${var_user}:${var_group} pack2.txt
154 chown ${var_user}:${var_group} hd.img
155 chown ${var_user}:${var_group} menu.lst

```

## **Part II.**

# **Custom Kernel for the x86 Architecture**

# 1. Booting a custom kernel

To boot a kernel in the x86 architecture the image must provide a global descriptor table and handle interrupt service routines adhering to certain arrangements.

## 1.1. Global Descriptor Table (GDT)

The x86 architecture has 6 different types of memory segments: code (*cs*), data (*ds*), stack (*ss*), extra (*es*), and two general purpose segments (*fs* and *gs*). For the limit purpose of this simple custom kernel, all segments except the *cs* will be set to point to the same portion of memory (ie 0x10), which is further explained in 1.1.3 below.

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Base address (bits 0-15)																Segment limit (bits 0-15)															
63	62	61	60	59	58	57	56	55	54	53	52	51	50	49	48	47	46	45	44	43	42	41	40	39	38	37	36	35	34	33	32
Base address (bits 24-31)								G	D	0	A	Limit (bits 16-19)				P	DPL		DT	T				Base address (bits 16-23)							

Table 1.1.: GDT memory structure

The above memory structure can be implemented using the following C code structure:

```

1 struct gdt_entry_struct {
2     u16int limit_low;
3     u16int base_low;
4     u8int  base_middle;
5     u8int  access;
6     u8int  granularity;
7     u8int  base_high;
8 } __attribute__((packed));

```

Code snippet 1.1: GDT entry structure

Notice that since elements that are less than 8 bits wide, one has to combine the bits into a single variable and manage setting them manually using bitwise operations. The `__attribute__((packed))` tells the compiler (GCC in this case) not to rearrange the fields in the structure. This is important as we want the exact arrangement in memory as that's what the x86 architecture is expecting.

### 1.1.1. Granularity (bits 52-55)

Bit 52 is available for OS usage. Bit 53 is reserved to be zero always. Bit 54 describes the size of the segment limit whether it is 16-bit (when set to 0) or 32-bit (when set to 1). Bit 55, the granularity (G) bit, defines the multiplier of the limit register (segment limit); if set to 1 the multiplier is set to 4 Kilobytes (4096 bytes), otherwise the multiplier is 1 (not modified).

### 1.1.2. Access (bits 40-47)

The structure of the access flag byte is a bit signifying the presence (P, bit 47) or otherwise of a segment in memory, which is useful for virtual memory management. The next 2 bits for defining the privilege

level of the descriptor (DPL, bits 46-47), which can take values between 0 and 3; this flag is sometimes called the ring level. When the ring level is set to 0 the process is the most privileged and runs in kernel-mode; on the other hand, ring level 3 is the least privileged, which is typically how most processes run in user-mode. This is very important as it enforces protection of certain CPU instructions from being run in higher ring levels. The following bit is the descriptor type bit, which if set means the the segment is a code/data segment, if not then it is a system segment. The next 4 bits represent the type of a segment (T); the first bit (bit 43) describes whether the segment is data or code, if set to 1, then the segment is a code segment. Bit 41 sets whether the segment is read only or read/write (for data segments) and execute only or execute/read (for code segments). This is important as it provides protections from process being possibly being exploited.

### 1.1.3. Code

```

1 gdt_set_gate(0, 0, 0, 0, 0);
2 gdt_set_gate(1, 0, 0xFFFFFFFF, 0x9A, 0xCF);
3 gdt_set_gate(2, 0, 0xFFFFFFFF, 0x92, 0xCF);
4 gdt_set_gate(3, 0, 0xFFFFFFFF, 0xFA, 0xCF);
5 gdt_set_gate(4, 0, 0xFFFFFFFF, 0xF2, 0xCF);
6
7 // assembly function call
8 gdt_flush((u32int)&gdt_ptr);

```

Code snippet 1.2: GDT segment creation

The `gdt_set_gate()` function is a helper function that simplifies loading 5 code and data segments into memory using the x86 GDT structure presented earlier. It manipulates the parameters and places them in an array (`gdt_entries`). `0xFFFFFFFF` is hexadecimal for about 4 billion (or  $2^{32}$ ), which is the maximum memory size an x86 system can have due to the base address being 32 bits. However, the first GDT entry must be made of all zeros this is called the null segment; a system will not boot without this segment.

The “access” parameter is explained below for each segment (after converting from hexadecimal to binary):

P	DPL		DT	T			
47	46	45	44	43	42	41	40
1	0	0	1	1	0	1	0

Table 1.2.: Kernel-mode code segment ( $0x9A=1001\ 1010_2$ )

P	DPL		DT	T			
47	46	45	44	43	42	41	40
1	0	0	1	0	0	1	0

Table 1.3.: Kernel-mode data segment ( $0x92=1001\ 0010_2$ )

P	DPL		DT	T			
47	46	45	44	43	42	41	40
1	1	1	1	1	0	1	0

Table 1.4.: User-mode code segment ( $0xFA=1111\ 1010_2$ )



P	DPL		DT	T			
47	46	45	44	43	42	41	40
1	1	1	1	0	0	1	0

Table 1.5.: User-mode data segment (0xF2=1111 0010<sub>2</sub>)

As can be seen in the above tables bit 43 is changed to differentiate between data and code segments. While bits 45 and 46, are used here to separate the kernel from the user memory segments.

P	D		A
55	54	53	52
1	1	0	0

Table 1.6.: Granularity (0xCF & 0xF0=1100 0000<sub>2</sub>)

This means that the segments are 32-bit wide and the multiplier of the limit register is 4K.

```

1 [GLOBAL gdt_flush] ; Allows the C code to call gdt_flush()
2
3 gdt_flush:
4     mov eax, [esp+4]
5     lgdt [eax]      ; Load the new GDT pointer
6
7     mov ax, 0x10    ; 0x10 is the offset in the GDT to data segment
8     mov ds, ax      ; Load all data segment selectors
9     mov es, ax
10    mov fs, ax
11    mov gs, ax
12    mov ss, ax
13    jmp 0x08:.flush ; 0x08 is the offset to code segment
14 .flush:
15     ret

```

Code snippet 1.3: GDT memory loading

To load these segments one must tell the CPU where to find the GDT entries table in memory. This can be done using the `lgdt` instruction to load a pointer of the GDT to memory. Since each GDT entry is 8 bytes and the kernel-mode data segment is the third, its offset is 8 (null segment) + 8 (code segment) = 16 bytes, which is 0x10 in hexadecimal; so we set all our segments to point to it, except for the kernel-mode segment. Finally, we need to start executing our kernel-mode code segment, to do so we need to goto second segment we loaded (after 8 bytes; 0x08), using the `jmp` instruction.

## 1.2. Interrupt Descriptor Table (IDT)

Interrupts in the x86 architecture can be generated by the CPU when a fault or an exception occurs (eg division by zero) or by software using the `int` instruction. There are 256 possible values of interrupt signals, and all of these must be defined even if they are just set to null. If some of these are not defined (nor null), then the CPU will panic and reset, as it expects them to be available.

Exception/interrupt value		Exception/interrupt meaning
decimal	hexadecimal	
0	0x00	Division by zero
1	0x01	Debugger
2	0x02	Non maskable interrupt
3	0x03	Breakpoint
4	0x04	Overflow
5	0x05	Out of bounds
6	0x06	Invalid opcode
7	0x07	No coprocessor
8	0x08	Double fault
9	0x09	Coprocessor segment overrun
10	0x0A	Invalid task state segment (TSS)
11	0x0B	Segment not present
12	0x0C	Stack fault
13	0x0D	General protection fault
14	0x0E	Page fault
15	0x0F	Unknown interrupt
16	0x10	Coprocessor fault
17	0x11	Alignment check
18	0x12	Machine check
19-31	0x13-0x1F	Reserved

Table 1.7.: Interrupt values and their meanings, as defined by the x86 architecture

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Kernel segment selector																Interrupt address (bits 0-15)															
63	62	61	60	59	58	57	56	55	54	53	52	51	50	49	48	47	46	45	44	43	42	41	40	39	38	37	36	35	34	33	32
Interrupt address (bits 16-31)																P	DPL			Type				0	0	0	0	0	0	0	0

Table 1.8.: IDT memory structure

The above memory structure can be implemented using the following C code structure:

```

1 struct idt_entry_struct {
2     u16int base_lo
3     u16int sel;
4     u8int  always0;
5     u8int  flags;
6     u16int base_hi;
7 } __attribute__((packed));

```

Code snippet 1.4: IDT entry structure

### 1.2.1. Constants and Type (bits 32-44)

Bits 32-39 must always be set to zero. Bits 40-43 (Type) is the type of the interrupt and is typically set to 6 (0110<sub>2</sub>) or 14 (1110<sub>2</sub>), which mean they are of type interrupt gate entry. Bit 44 should be zero if bits 40-43 are of type interrupt gate. These constants are represented in the memory structure above.

### 1.2.2. DPL and P (bits 45-47)

The DPL bits define the minimum required privilege to be able to handle the interrupt. This makes sure certain interrupts are protected from user-space. On the other hand the P bit defines whether the interrupt is present (ie whether the interrupt is handled or not). If this bit is set to zero for a certain interrupt signal, then if this interrupt occurs the CPU generates an “interrupt not handled” exception.

### 1.2.3. Selector (bits 16-32)

These bits are the relative address in the GDT table of the kernel-mode code segment, which is 0x08 (as explained above 1.1.3).

### 1.2.4. Code

```
1 | idt_set_gate(0, (u32int) isr0, 0x08, 0x8E);
2 |
3 | // assembly function call
4 | idt_flush((u32int)&idt_ptr);
```

Code snippet 1.5: IDT handler memory locations

The `idt_set_gate()` function is a helper function that simplifies loading of interrupt handlers into using the x86 IDT structure presented earlier. It manipulates the parameters and places them in an array (`idt_entries`). The first parameter is the index of the interrupt in the array; in other words the exception value that will be handled by the next parameter. The next parameter is a pointer to the interrupt service routine handler we want to run upon receiving a certain interrupt. The following parameter is the kernel segment selector. Finally the last parameter is the “flags”, which is represented by bits 40-47. These bits are inspected further below:

P	DPL		Type				
47	46	45	44	43	42	41	40
1	0	0	0	1	1	1	0

Table 1.9.: IDT flags (0x8E=1000 1110<sub>2</sub>)

Bit 47 (P) is set because we are handling this exception. The DPL bits (45-46) are set to zero, which means to handle these exception a segment must have a privilege ring level 0 (ie only kernel-mode segments can handle these exceptions). Finally the type is set to 14 (1110<sub>2</sub>), which as described earlier means this is an interrupt gate handler.

```
1 | [GLOBAL idt_flush] ; Allows the C code to call idt_flush()
2 |
3 | idt_flush:
4 |     mov eax, [esp+4]
5 |     lidt [eax]      ; Load the IDT pointer
6 |     ret
```

Code snippet 1.6: IDT memory loading

To load these interrupt service routines one must tell the CPU where to find the IDT entries table in memory. This can be done using the `lidt` instruction to load a pointer of the IDT to memory. Afterwards, the CPU knows that exceptions from 0 to 31 are handled and should call each handler respectively.

## 2. Monitor

GRUB boots the kernel in text mode, which means the kernel can directly access the framebuffer of the screen, as a set of characters directly mapped to memory. The screen size provided by GRUB is 80 characters wide by 24 characters long. It can be accessed by writing directly to the 0xB8000 address in memory that is mapped to the video memory of the VGA controller, which is not necessarily the RAM. The VGA has two important registers the control (0x3D4) and data (0x3D5) registers that can be used to update the location of the cursor. The framebuffer is an array of 16-bit elements using the following structure for each element:

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Background color				Foreground color				ASCII character code							

Table 2.1.: Framebuffer element structure

The color values are 8 bits wide each, as seen in the table above. which means that only 16 colors can be displayed. This leaves another 8 bits for the character code, as such only ASCII characters can be displayed using this method.

Each character on the screen can be accessed using the following formula:

$$character\ address(x, y) = B8000_{16} + (80_{10} \cdot y + x)$$

For example, if we want to write a character at column 2 ( $y = 3$ ) and row 4 ( $x = 5$ ); we will have to write the 16-bit character element to address  $0xB80F5_{16}$ .

### 3. Debugging the kernel

The easiest method of debugging this minimal kernel is by using tools such as `gdb`. As GRUB hands control to the system to our kernel, it is useful to set breakpoints at `gdt_flush` and `idt_flush` because that way one could inspect the values and contents of the GDT and the IDT elements being loaded into memory. After initialization of the GDT and IDT, setting breakpoints at `monitor_write` and `isr_handler`. Setting a breakpoint at `monitor_write` helps in knowing what is being printed to the screen and when. Moreover, setting a breakpoint at `isr_handler`, helps in reading the interrupt values being received and that will be printed using the monitor set of functions.